



Linux Academy

Ansible Setup, Configuration, and Ad-Hoc Commands Deep Dive

Course Outline Companion

Stosh Oldham

stosh@linuxacademy.com

January 30, 2019

Contents

How to Install Ansible	1
Configure SSH users (with sudo) and keys	1
The Ansible Configuration File	2
Setting Up the Ansible inventory	3
The Ansible Command	4
Understanding Ansible Modules	4
The Shell and Command modules	5
Collecting System information	5
Working with the File and Copy Modules	6
Editing File Contents with the <code>lineinfile</code> Module	7
Downloading Files with the <code>get_url</code> Module	7
Working with File Archives	8
Creating System Users with the User Module	9
Working with the Group Module	9
Installing Software	9
Controlling Daemons with the Service Module	10
Managing Long-running Commands	10
Parallelism in Ansible Ad-hoc	11

How to Install Ansible

- In order to install Ansible, you must configure the EPEL repository on your system.
- Once the EPEL repository is configured, your package manager can install Ansible and manage dependencies.
 - `sudo yum install ansible`
- It is advisable to install some means of source control as well.
 - `sudo yum install git`
- After package installation, there are two further steps to be taken before Ansible may be used.
 - SSH access must be provided on host you wish to control with Ansible.
 - The Ansible inventory must be configured for each host you wish to control with Ansible.
 - These points are covered in the next sections.

Configure SSH users (with sudo) and keys

- Ansible is best implemented using a common user across all Ansible controlled systems.
- Creating a System User for Ansible
 - Only a basic system user with ssh access is needed for Ansible to connect to a host.
 - The following commands, executed as root, will create a new user on a system called **ansible** and allow for the user password to be set:
 - `useradd ansible`
 - `passwd ansible`
 - This step should be performed on the Ansible control node and repeated on each other node you wish to have Ansible control.
- Pre-shared Key Authentication
 - While it is possible to connect to a remote host with Ansible using password authentication using `-k` (note lowercase), it is not a common practice as it can incur significant overhead in terms of manual intervention.
 - The `ssh-keygen` and `ssh-copy-id` command can facilitate creating a pre-shared key for user authentication.
 - As the **ansible** user, first run `ssh-keygen` without parameters and follow the prompts to create a key pair for the user.
 - It is recommended that you perform this step on your Ansible control node.
 - Once the key pair has been created, run `ssh-copy-id ansible@TARGET_HOST` to copy the **ansible** user's public key to TARGET_HOST.

- (**Note:** The `ansible` user must already exist on `TARGET_HOST` and be able to log in with password authentication.)
- After the public key is in place, you should be able to connect to `TARGET_HOST` as the `ansible` user from your original host without being prompted for a password.
- Escalating permissions
 - For effective system management, the `ansible` user will need a means of privilege escalation.
 - `/etc/sudoers` may be edited to allow your selected user to sudo any command without password for the most automated configuration using the line `ansible ALL=(ALL) NOPASSWD: ALL`.
 - It is also possible to prompt for a sudo password at runtime using `-K`(note uppercase) if desired (**Note:** This can become a challenge when executing against many systems).

The Ansible Configuration File

- Ansible has a number of settings that may be adjusted.
- The default Ansible configuration file is `/etc/ansible/ansible.cfg`
- Notable configurations include:
 - Default inventory configuration
 - Default remote user
- The configuration properties may be overridden using local files.
 - The first configuration file found is used and later files are ignored.
 - Configuration file which will be searched for is in the following order:
 - `ANSIBLE_CONFIG` (environment variable if set)
 - `ansible.cfg` (in the current directory)
 - `~/.ansible.cfg` (in the home directory)
 - `/etc/ansible/ansible.cfg`
- Starting at Ansible 2.4, the `ansible-config` utility ships with Ansible.
 - The utility allows users to see:
 - All the configuration settings available
 - Their defaults
 - How to set them
 - Where their current value comes from.
- The utility may use a few different sub-commands:
 - `view`: Displays the current config file.
 - `list`: Provides all current configs reading `lib/constants.py` and shows env and config file setting names.

- **dump**: Shows the current settings, merges `ansible.cfg` if specified.
- **Note**: the **dump** sub-command may take the flag `--only-changed` which will only show configurations not set to the default.
- The `-c` or `--config` switches may be used to specify a particular configuration file.
- If one is not provided, the highest precedent configuration file is used.
- See `man ansible-config` for more information.

Setting Up the Ansible inventory

- An inventory is a list of hosts that Ansible manages.
- The default Ansible Inventory File is `/etc/ansible/hosts`.
- Inventory location may be specified as follows:
 - Default: `/etc/ansible/hosts`
 - Specified by CLI: `ansible -i <filename>`
 - Can be set in `ansible.cfg`
- Example Inventory file:

```
mail.example.com ansible_port=5556 ansible_host=192.168.0.10
```

```
[webservers]
httpd1.example.com
httpd2.example.com
```

```
[labservers]
lab[01:99]
```

- The first line of the example file defines a host, `mail.example.com`.
- Two variables are affiliated with the host, `ansible_port` and `ansible_host`.
- The groups `webservers` and `labservers` are defined in this example.
- Note the `labservers` group has 99 hosts in it that are defined via a pattern
- (The expression `lab[01:03]` is the same as specifying `lab01`, `lab02`, `lab03`.)
- Note that inventory files may be provided in YAML format but that is uncommon.
- It is recommended that you keep separate inventories for staging and production.

The Ansible Command

- Ansible ad-hoc commands analogous to bash commands.
- Playbook analogous to bash script (but, that is another course).
- Syntax: `ansible <HOST> -b -m <MODULE> -a "<ARG1 ARG2 ARGN>" -f <NUM_FORKS>`
 - **HOST** is a host or host group defined in the Ansible inventory file.
 - **-b** is for 'become' (This replaces the depreciated **-s** flag as in "sudo").
 - Ansible escalates permission to **--become-user** using the method defined by **--become-method**.
 - Default become-user is **root**.
 - Default become-method is **sudo**.
 - **-m** is for module or command to use (We will look at modules in the next section).
 - **-a** is for parameters to pass. If used without module, it is like running a shell command on the target system(s).
 - **-f** is used to set forks for parallelism, which is how you can have Ansible execute plays simultaneously on many hosts.

Understanding Ansible Modules

- Modules are discrete units of code that can be used from the command line or in a playbook task.
 - Modules may take arguments.
 - Most modules take **key=value** arguments, delimited by whitespace.
 - There are a couple exceptions found in the command and shell modules which we will get to later.
- Common modules (more detail on each and more modules later):
 - Module: **ping**
Purpose: Verify Ansible connectivity between hosts.
Arguments: None
 - Module: **setup**
Purpose: Collect Ansible facts.
Arguments: None
 - Module: **yum**
Purpose: Install software with yum.
Arguments: **name=<PACKAGE_NAME> state=<STATE>**
 - Module: **service**
Purpose: Control service daemons.
Arguments: **name=<SERVICE_NAME> state=<STATE>**
 - Module: **copy**
Purpose: Copy a file to a particular location on a target host.
Arguments: **src=<SOURCE_PATH> dest=<ABSOLUTE_DESTINATION_PATH>**

- When you need help
 - There are a very large number of modules that perform a variety of tasks.
 - It is nearly impossible to know every module.
 - That said, there is good documentation provided on *docs.ansible.com*.
 - The module index is provided at *docs.ansible.com* and it provides detailed information on each module.
 - Ansible ships with the **ansible-doc** command as well.
 - Specifying a module name as a parameter will provide module specific documentation.
 - The **-l** flag will list installed modules with a brief description.

The Shell and Command modules

- The **shell** and **command** modules are both effective ways to run 'raw' commands on a target host.
- The **command** module
 - It is the default module used with Ansible ad-hoc when no module is specified.
 - The arguments provided are 'free form' and ran as a command on each specified remote host.
 - The module itself may take a few parameters:
 - **creates**: Execution of the module results in the creation of the provided file.
 - **removes**: Execution of the module results in the removal of the provided file.
 - **chdir**: Change to the provided directory before executing the command.
- The **shell** modules
 - Works almost exactly like the command module except commands are ran in a shell environment (using **/bin/sh**).
 - The primary reason the **shell** module might be used over the **command** module is for use of input / output redirection, environment variables, or other such shell features (which are not supported with the **command** module).
 - Per the Ansible documentation, the best practice is to use the command module unless you require shell features.
 - **Pro tip**: If running multiple shell commands is necessary, there is a **script** module that will copy a shell script to a remote host and execute it, featuring the same arguments that are supported by the **shell** command. (Alternatively, you can create your own Ansible module!)

Collecting System information

- Ansible facts are simply various properties regarding a given remote system.
- The **setup** module can retrieve facts.
 - The filter parameter takes regex to allow you to prune fact output.

- Facts are gathered by default in Ansible playbook execution.
 - The keyword **gather_facts** may be set in a playbook to change fact gathering behavior.
- It is possible to print Ansible facts in files using variables.
- Facts may be filtered using the **setup** module ad-hoc by passing a value for the **filter** parameter.
- Ansible command output may be directed to a file using the **--tree outfile** flag, which may be helpful when working with facts.
- It is possible to use **{{ ansible_facts }}** for conditional plays based on facts.

Working with the File and Copy Modules

- Working with files is an essential operation for systems management and configuration.
- Ansible has a number of core modules that provide for file manipulation that will be covered throughout this course.
- The **file** module
 - The **file** module can be used to create, delete, and modify file properties.
- The **copy** module
 - This module is used for copying files.
 - Files may be copied from a number of sources.
 - Files may be copied from the control node to target nodes.
 - Files may be copied from file systems on the target node.
 - Files may be created during execution of the copy module and place on target nodes.
- The following are notable arguments used by the **copy** module:
 - **remote_src**
 - If no, it will search for **src** on the control node.
 - If yes it will go to the target node for the **src**.
 - Default is no.
 - Currently **remote_src** does not support recursive copying.
 - **remote_src** only works with **mode=preserve** as of version 2.6.
 - **content**: When used instead of **src**, sets the contents of a file directly to the specified value.
 - **dest**: Remote absolute path where the file should be copied to.
 - **mode** - Mode (in octal or symbolic notation) the file or directory should be.
 - **src**
 - Local path to a file to copy to the remote server; can be absolute or relative.

- If path is a directory, it is copied recursively.
 - In this case, if path ends with `"/"`, only contents inside of that directory are copied to destination.
 - Otherwise, if it does not end with `"/"`, the directory itself with all contents is copied.
- This behavior is similar to `rsync`.

Editing File Contents with the `lineinfile` Module

- The `lineinfile` module is a key way of interacting with a file's contents.
- The key feature of the module is inserting a line of provided text into a given file.
- The module will place the line of text at the end of a file by default, however it can place the line after or before a line of matched text.
- If the module detects the provided text already in place, it will not add it again.
- It may also remove lines of text.
- Key arguments for the `lineinfile` module:
 - **backup**: Create a backup file, including the timestamp information, so you can get the original file back if you somehow clobbered it incorrectly.
 - **create**: Create the file if it does not exist. (False by default.)
 - **state**: Whether or not the line should be there.
 - **insertafter**: Takes a regular expression and inserts `line` *after* the first matched line. (Only honored if **regexp** is undefined or unmatched.)
 - **insertbefore**: Takes a regular expression and inserts `line` *before* the first matched line. (Only honored if **regexp** is undefined or unmatched.)
 - **regexp**: Takes a regular expression and replaces the last line matched with **line** or removes the matched line if **state** is set to *absent*.
 - **line**: The text to insert/replace.
 - **path**: The file to modify.
- **Note**: There is a `replace` module that can make more granular file changes using regular expressions.
 - It uses a **replace** argument instead of the **line** argument.
 - The module also supports **before**, **after**, **path**, and **regexp** arguments.

Downloading Files with the `get_url` Module

- It is frequently necessary to download files from HTTP hosts on managed hosts.
- The `get_url` module provides functionality for downloaded files over HTTP with a variety of options.
 - The HTTP, HTTPS, and FTP protocols are supported

- Basic and SSL Authentication
- Checksum verification
- Proxy network access
- Notable Arguments:
 - **dest**: Where the file should be placed on the *remote host*.
 - **checksum** - The algorithm and result value used to verify the file's checksum.
 - **Note**: The result value may be provided as a text file containing the value or the plain text value.
 - **force_basic_auth**: This option forces the sending of the Basic authentication header upon initial request.
 - **mode**: The mode of the final file.
 - **url**: The URL of the file to download.

Working with File Archives

- The **archive** and **unarchive** modules may be used to to work with various types of file archives.
- These modules work with common archives such as **tar**, **gzip** (default), and **zip**.
- These modules allow for compression (or expansion) of files and directories.
- **archive** Notable Arguments:
 - **dest**: The file name of the archive file.
 - **format**: Select compression format.
 - **path**: Remote absolute path, glob, or list of paths or globs for the files to archive.
 - **remove**: Controls if the source files will be deleted after the archive is created.
- **unarchive** notable arguments:
 - **dest**: Where archive files should be unpacked.
 - **remote_src**: If set to 'yes', file is already on target system.
 - **remove**: Controls if the source archive file will be deleted after execution.
 - **src**: Location of source archive.
 - If **remote_src** is set to 'no' (the default) **src** is local to control server.
 - If **remote_src** is set to 'yes', **src** is on target server.
 - If **remote_src** is set to 'yes' and **src** contains **://**, the target server will download the file from specified URL (similar to **get_url** module).

Creating System Users with the User Module

- Interacting with system users is a common operational task.
- Ansible allows this functionality using the **user** module.
 - The **user** module may interact with users in many standard ways.
 - Create users
 - Remove users
 - Change user properties
- Common Arguments:
 - **group**: Specify a primary group for the named user.
 - **groups**: Specify secondary groups for the named user.
 - **name**: Name of the user with which to interact.
 - **generate_ssh_key**: Whether to generate a SSH key for the named user.
 - **remove**: When used with **state=absent**, the named user is deleted.
 - **state**: Creates the named user if it doesn't exist when set to 'present'.
 - **uid**: Optionally set the UID of the named user.

Working with the Group Module

- Just as with system users, Ansible support working with system groups.
- The **group** module serves this function.
 - The **group** module works with groups in the same way the **user** module interacts with users.
- Common Arguments:
 - **gid**: Optionally set the GID of the named group.
 - **name**: Name of the group with which to interact.
 - **state**: Creates the named group if it doesn't exist when set to 'present'.
 - **system**: Indicate if a group is a system group.

Installing Software

- One of the key features necessary for bootstrapping and managing a host is the ability to install and manage software packages.
- **package**, **yum**, and **apt** are used for this purpose depending on distribution.
 - **yum** and **apt** are used for distribution specific software management.

- The **package** module will automatically detect the target host distribution and use the appropriate method of installing software.
- In a mixed distribution environment, the **package** module is best practice.
- Common Arguments:
 - **name**: The name of the software package to be installed.
 - **state**: If the named package should be 'present' or 'absent'.
 - **user**: Which package manager to use. (Defaults to 'auto'.)

Controlling Daemons with the Service Module

- Another common administration task is to control system daemons.
- The **service** module serves this purpose.
 - It is compatible with both systemd and the older BSD init style of daemon management.
- Common Arguments:
 - **enabled**: Configures whether or not the service should start on boot.
 - **name**: The name of the service to manage.
 - **state**: The state of the service after the module executes. May be one of:
 - *started*: Service will be started (will not restart if already started).
 - *stopped*: Service will be stopped.
 - *restarted*: Service will be restarted even if it is already started.
 - *reloaded*: Service configuration will be reloaded if supported by the service.

Managing Long-running Commands

- Some operations may require a significant amount of time to execute.
- Ansible provides a feature to allow an operation to run asynchronously such that the status may be checked.
- There are two options related to running asynchronous commands:
- **-B** is used to provide a timeout value and initiate the operation.
 - Unit: seconds
 - Default value: N/A
 - Job is terminated when timeout value expires.
- **-P** may be provided to engage polling of the operation on a set interval.
 - Unit: seconds

- Default value: 15
- Polling will leave the Ansible command in the foreground until execution on all hosts completes.
- Set to 0 to disable polling.
- Example: `ansible all -B 1200 -P 60 -a "/usr/bin/long_running_operation"`
- This example provides a 20 minute timeout and polls the operation's status every 60 seconds.
- The status may be checked using the module `async_status`.
 - Example: `ansible httpd1 -m async_status -a "jid=12334568901234.1234"`
 - Note: JID is only realistically obtained during playbook operation.

Parallelism in Ansible Ad-hoc

- Ansible uses what are known as forks to execute tasks in parallel.
- Each fork is able to run a task against a target host.
- By default, Ansible executes 5 forks.
- Assuming one runs a module against 5 hosts, the module should execute near simultaneously on all hosts.
- If more than 5 hosts are targeted, hosts after the first 5 must wait for a fork to become free before being able to execute.
- The number of forks used by Ansible may be configured in `ansible.cfg` or on an as-needed basis using the `-f` or `--forks` flag.
- **Note:** The more forks allowed, the more system resources will be required on the Ansible control node.
- Example: `ansible all -f 10 -m setup`
 - This example will use 10 forks to execute the `setup` module against all hosts in the inventory.
 - **Note:** Ansible will only fork as much as required.
 - Presuming Ansible is configured to run 10 forks but is only tasked with executing on 6 hosts, only 6 forks will be created.