



Linux Academy

Ansible Quick Start

Just the facts...

Stosh Oldham

stosh@linuxacademy.com

May 29, 2019

Contents

- Install and Configure Ansible 1
- When you need help 2
- Ansible ad-hoc 2
- Ansible Playbooks 3
- Ansible Variables 4
- Ansible Facts 5
- Debugging in Ansible 6
- Notifications and Handlers 6

Install and Configure Ansible

- In order to install Ansible, you must configure the EPEL repository on your system.
- Once the EPEL repository is configured, your package manager installs Ansible and manage dependencies.
 - `sudo yum install ansible`
- It is advisable to install some means of source control as well.
 - `sudo yum install git`
- Configuration Files:
 - `/etc/ansible/ansible.cfg`:
 - The primary Ansible configuration file.
 - Notable configurations include:
 - Default inventory configuration.
 - Default remote user.
 - `/etc/ansible/hosts`:
 - Default Ansible Inventory File.
 - An inventory is a list of hosts that Ansible manages.
 - Inventory location may be specified as follows:
 - Default: `/etc/ansible/hosts`.
 - Specified by CLI: `ansible -i <filename>`.
 - Can be set in `ansible.cfg`.
 - Example Inventory file:

```
mail.example.com ansible_port=5556 ansible_host=192.168.0.10

[webservers]
httpd1.example.com
httpd2.example.com

[labservers]
lab[01:99]
```

 - The first line defines a host `mail.example.com`.
 - Two variables are affiliated with the host, `ansible_port` and `ansible_host`.
 - The group's web servers and lab servers are defined in this example.
 - Note the lab servers group has 99 hosts in it that are defined via a pattern.
 - The expression `lab[01:03]` is the same as specifying `lab01`, `lab02`, `lab03`.
 - Configure SSH users (with `sudo`) and keys:
 - While it is possible to connect to a remote host with Ansible using password authentication using `-k` (note lowercase), it is not a common practice as it can incur significant overhead concerning manual intervention.

- Ansible is best implemented using a common user across all Ansible controlled systems.
- The `ssh-keygen` and `ssh-copy-id` command can facilitate creating a pre-shared key for user authentication.
- `/etc/sudoers` may be edited to allow your selected user to `sudo` any command without a password for the most automated configuration using the line `ansible ALL=(ALL) NOPASSWD: ALL`.
- It is also possible to prompt for a `sudo` password at runtime using `-K`(note uppercase) if desired; this can become a challenge when executing against many systems.

When you need help

- Useful documentation is provided at docs.ansible.com.
- A module index is provided at docs.ansible.com that provides detailed information on each module.
- Ansible ships with the `ansible-doc` command which:
 - Specifies a module name as a parameter will provide module specific documentation.
 - Uses the `-l` flag to list installed modules with a brief description.

Ansible ad-hoc

- Ansible ad-hoc commands are analogous to bash commands.
- Playbooks are analogous to bash scripts.
- Syntax: `ansible \<HOST> -b -m \<MODULE> -a "\<ARG1 ARG2 ARGN>" -f \<NUM_FORKS>:`
 - `HOST` is a host or host group defined in the Ansible inventory file.
 - `b` is for become:
 - Replaces the deprecated `-s` flag as in `sudo`.
 - Ansible escalates permission to `--become-user` using the method defined by `--become-method`.
 - Default `become-user` is `root`.
 - Default `become-method` is `sudo`.
 - `m` is for modules for the command to use.
 - `a` is used for parameters to pass, if used without `m`, it is like running a shell command on the target system(s).
 - `f` is used to set forks for parallelism, which is how you can have Ansible execute plays simultaneously on many hosts.
- Common modules with required parameters:
 - `ping`
 - `setup`
 - `yum "name=\<NAME> state=\<STATE>"`
 - `service "name=\<NAME> state=\<STATE>"`
 - `copy "src=\<SOURCE_PATH> dest=\<ABSOLUTE_DESTINATION_PATH>"`

Ansible Playbooks

- Basic Ansible Playbook structure:
 - As ad-hoc commands are to bash commands, playbooks are to bash scripts.
 - Playbooks are ran using the `ansible-playbook` command, not the `ansible` command.
 - You must specify the playbook to run as a parameter.
 - Playbooks are written in YAML:
 - Each play is an element in a sequence.
 - Plays contain lists of hosts and, at minimum, one or more tasks.
 - Each task has a name and module.
 - Modules may have parameters.
 - Improper indentation can cause a playbook to err in a vague way.
 - Spaces matter.

- Sample playbook:

```
- hosts: webservers
  become: yes

  tasks:
    - name: ensure apache is at the latest version
      yum:
        name: httpd
        state: latest
    - name: write the apache config file
      template:
        src: /srv/httpd.j2
        dest: /etc/httpd.conf

- hosts: databases
  remote_user: root

  tasks:
    - name: ensure postgresql is at the latest version
      yum:
        name: postgresql
        state: latest
    - name: ensure that postgresql is started
      service:
        name: postgresql
        state: started
```

- The above playbook is composed of two plays:
 - The first play installs a package (using the `yum` module) and creates a configuration file from a template (using the `template` module):

- Note that `become` is set for the first play:
 - This is the same as using `-b` when running an ad-hoc command.
 - Ansible will use `sudo` to escalate to the `root` account on the remote system.
 - This is required for the tasks defined.
- The second play ensures `postgresql` is installed and running using both the `yum` and `service` modules.
- Note that `become` is *not* set in the second play yet the `remote_user` is set to `root`:
 - This is not a particularly good security practice but does demonstrate some basic Ansible configuration.
- Check Mode provides a quick sanity check:
 - Dry run that does not affect the systems.
 - Syntax: `ansible-playbook foo.yml --check`
- Retry file:
 - If a playbook fails, a retry file is generated containing the list of hosts where the play failed
 - A file called `\<playbook-name>.retry` is created upon a playbook failure.
 - The file may be specified using `--limit` with the same playbook to reattempt the playbook at a later time.
- Plays should be safe to run repeatedly against the same target without ill effect.

Ansible Variables

- Variable names should be letters, numbers, and underscores.
- Variables should always start with a letter.
- Examples of valid variable names:
 - `foobar`
 - `foo_bar`
 - `foo5`
- Examples of invalid variable names:
 - `foo-bar`
 - `1foobar`
 - `foo.bar`
- Variables can be scoped by group, host, or within a playbook.
- Variables are passed in via the command line using the `--extra-vars` or `-e` flag or are defined within a playbook:

- CLI Example: `ansible-playbook service.yml -e "target_hosts=localhost target_service=httpd"`
- Playbook Example:

```
hosts: webservers
become: yes
vars:
  target_service: httpd
  target_state: started
tasks:
  - name: Ensure target service is at target state
    service:
      name: "{{ target_service }}"
      state: "{{ target_state }}"
```

- Variables are referenced using double curly braces.
- It is good practice to wrap variable names or statements containing variable names in weak quotes:
 - Example: `hosts: "{{ my_host_var }}"`
- Typical uses of variables:
- Customize configuration values.
- Hold return values of commands.
- Ansible has many default variables for controlling Ansible's behavior.

Ansible Facts

- Ansible facts are simply various properties regarding a given remote system.
- The `setup` module can retrieve facts:
 - The `filter` parameter takes regex to allow you to prune fact output.
- Facts are gathered by default in Ansible playbook execution:
 - The keyword `gather_facts` may be set in a playbook to change fact gathering behavior.
- It is possible to print Ansible facts in files using variables.
- Facts may be filtered using the `setup` module ad-hoc by passing a value for the `filter` parameter.
- Ansible command output may be directed to a file using the `--tree outputfile` flag which may be helpful when working with facts.
- It is possible to use `{{ ansible_facts }}` for conditional plays based on facts.

Debugging in Ansible

- The debug module may be used to help troubleshoot plays:
 - Used to print detail information about in-progress plays.
 - Handy for troubleshooting.
- Debug takes two primary parameters that are mutually exclusive:
 - `msg`: A message that is printed to STDOUT
 - `var`: A variable whose content is printed to STDOUT.
 - Example:
 - `debug:`
`msg: "System {{ inventory_hostname }} has uuid {{ ansible_product_uuid }}"`
- The register module is used to store task output:
 - It essentially can save results of a command.
 - Several attributes are available: return code, stderr, and stdout.
 - Example:
 - `hosts: all`
`tasks:`
 - `shell: cat /etc/motd`
`register: motd_contents`
 - `shell: echo "motd contains the word hi"`
`when: motd_contents.stdout.find('hi') != -1`

Notifications and Handlers

- Ansible provides a mechanism that allows an action to be flagged for execution when a task performs a change.
- By only executing certain tasks during a change, plays are more efficient.
- This mechanism is known as a *handler* in Ansible.
- A handler may be called using the `notify` keyword.
- No matter how many times a handler is flagged in a play, it is only ran once during a play's final phase.
- `notify` will only flag a handler if a task block makes changes:
 - Example:


```
- name: template configuration file
  template:
    src: template.j2
    dest: /etc/foo.conf
  notify:
    - restart memcached
    - restart apache
```

- The calls made in the notify section correspond to handler definitions within the play.
- A handler may be defined similarly to tasks:

- Example:

```
handlers:
- name: restart cache service
  service:
    name: memcached
    state: restarted
    listen: "restart memcached"
- name: restart web services
  service:
    name: apache
    state: restarted
    listen: "restart apache"
```